

GPU Pareto Frontier: Accuracy vs. Peak Activation Memory across Batch Sizes

